```
In [1]: import tensorflow as tf
         print("TensorFlow version:", tf.__version__)
       TensorFlow version: 2.19.0
In [1]: import matplotlib.pyplot as plt
         print("  matplotlib working!")
        matplotlib working!
In [2]: import numpy as np
         import matplotlib.pyplot as plt
         import tensorflow as tf
         from tensorflow.keras.models import Sequential
         from tensorflow.keras.layers import Dense, Flatten
         from tensorflow.keras.datasets import mnist
         from tensorflow.keras.utils import to categorical
In [3]: # Load dataset
         (X_train, y_train), (X_test, y_test) = mnist.load_data()
         # Normalize image pixel values
         X_{train} = X_{train} / 255.0
         X_{\text{test}} = X_{\text{test}} / 255.0
         # One-hot encode the labels
         y_train_cat = to_categorical(y_train, num_classes=10)
        y test cat = to categorical(y test, num classes=10)
       Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mn
       ist.npz
       11490434/11490434 —
                                        3s 0us/step
In [4]: model = Sequential([
                                                    # Flatten 28x28 to 784
             Flatten(input_shape=(28, 28)),
             Dense(128, activation='relu'), # Hidden layer 1
Dense(64, activation='relu'), # Hidden layer 2
Dense(10, activation='softmax') # Output layer for
             Dense(10, activation='softmax')
                                                    # Output layer for 10 digits
         ])
         # Compile the model
         model.compile(optimizer='adam',
                       loss='categorical_crossentropy',
                       metrics=['accuracy'])
         # Show model summary
         model.summary()
       C:\Users\shouvik\anaconda3\envs\tf-env\lib\site-packages\keras\src\layers\reshaping
       \flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a
       layer. When using Sequential models, prefer using an `Input(shape)` object as the fi
       rst layer in the model instead.
         super(). init (**kwargs)
      Model: "sequential"
```

Layer (type)	Output Shape	
flatten (Flatten)	(None, 784)	
dense (Dense)	(None, 128)	
dense_1 (Dense)	(None, 64)	
dense_2 (Dense)	(None, 10)	

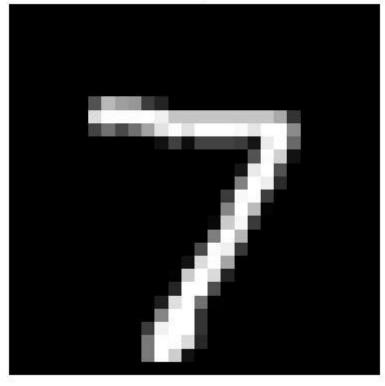
```
Total params: 109,386 (427.29 KB)
      Trainable params: 109,386 (427.29 KB)
      Non-trainable params: 0 (0.00 B)
In [5]: history = model.fit(X_train, y_train_cat, epochs=10, validation_split=0.2)
      Epoch 1/10
                         5s 3ms/step - accuracy: 0.8624 - loss: 0.4720 - val_a
      1500/1500 -
      ccuracy: 0.9604 - val loss: 0.1338
      Epoch 2/10
      1500/1500 —
                                  — 4s 3ms/step - accuracy: 0.9657 - loss: 0.1141 - val_a
      ccuracy: 0.9677 - val loss: 0.1079
      Epoch 3/10
                                  — 4s 3ms/step - accuracy: 0.9749 - loss: 0.0810 - val_a
      1500/1500 -
      ccuracy: 0.9666 - val loss: 0.1078
      Epoch 4/10
                                --- 5s 3ms/step - accuracy: 0.9829 - loss: 0.0560 - val_a
      1500/1500 —
      ccuracy: 0.9707 - val loss: 0.0991
      Epoch 5/10
      1500/1500 -
                                  4s 3ms/step - accuracy: 0.9877 - loss: 0.0409 - val a
      ccuracy: 0.9757 - val loss: 0.0830
      Epoch 6/10
                          4s 3ms/step - accuracy: 0.9901 - loss: 0.0303 - val_a
      1500/1500 —
      ccuracy: 0.9756 - val_loss: 0.0861
      Epoch 7/10
      1500/1500 5s 3ms/step - accuracy: 0.9903 - loss: 0.0288 - val_a
      ccuracy: 0.9710 - val_loss: 0.1096
      Epoch 8/10
      1500/1500 -
                                 — 4s 3ms/step - accuracy: 0.9932 - loss: 0.0218 - val_a
      ccuracy: 0.9722 - val loss: 0.1164
      Epoch 9/10
      1500/1500 -
                                 — 4s 3ms/step - accuracy: 0.9931 - loss: 0.0220 - val_a
      ccuracy: 0.9718 - val_loss: 0.1193
      Epoch 10/10
                                  - 4s 3ms/step - accuracy: 0.9938 - loss: 0.0174 - val a
      1500/1500 -
      ccuracy: 0.9762 - val_loss: 0.1064
In [6]: test_loss, test_acc = model.evaluate(X_test, y_test_cat)
        print(f"Test Accuracy: {test_acc:.2f}")
      313/313 — 1s 2ms/step - accuracy: 0.9722 - loss: 0.1114
      Test Accuracy: 0.98
```

```
In [7]: # Predict probabilities for test set
predictions = model.predict(X_test)

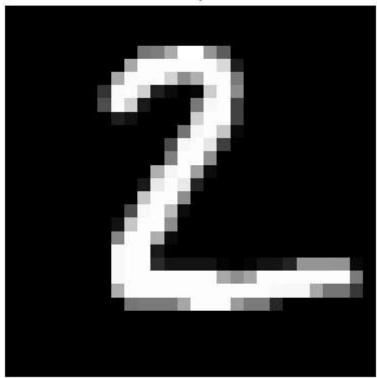
# Show 5 sample predictions
for i in range(5):
    plt.imshow(X_test[i], cmap='gray')
    plt.title(f"Predicted: {np.argmax(predictions[i])}, Actual: {y_test[i]}")
    plt.axis('off')
    plt.show()
```

313/313 Os 1ms/step

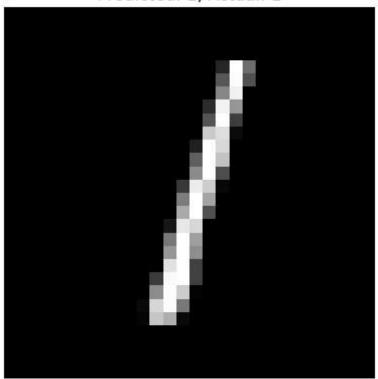
Predicted: 7, Actual: 7



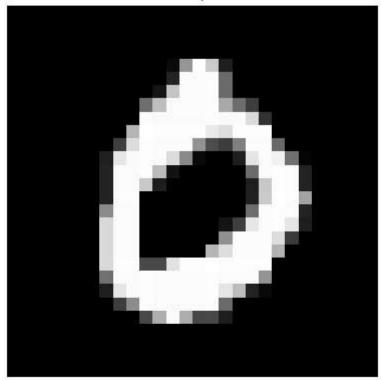
Predicted: 2, Actual: 2



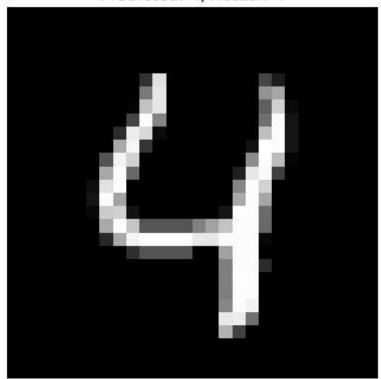
Predicted: 1, Actual: 1



Predicted: 0, Actual: 0



Predicted: 4, Actual: 4



Task 4 Summary – Handwritten Digit Classifier (MNIST)

© Objective:

To build a neural network that classifies handwritten digits (0–9) using the MNIST dataset. The model was built using TensorFlow/Keras Sequential API and evaluated on test data for accuracy and prediction performance.

Dataset Used:

- MNIST dataset:
 - 60,000 training images
 - 10,000 test images
 - Each image: 28×28 grayscale
 - Labels: Digits from 0 to 9

Steps Performed:

- 1. Data Loading & Normalization
 - Loaded data using tensorflow.keras.datasets.mnist
 - Scaled pixel values to [0, 1]
- 2. Label Encoding
 - Applied one-hot encoding to the labels using to_categorical()
- 3. Model Architecture (Sequential):
 - Flatten input layer (28×28 → 784)
 - Dense (128 units, ReLU)
 - Dense (64 units, ReLU)
 - Dense (10 units, Softmax)
- 4. Model Compilation & Training
 - Optimizer: Adam
 - Loss: Categorical Crossentropy
 - Epochs: 10
 - Validation Split: 20%
- 5. Evaluation & Accuracy
 - Evaluated model on test set
 - Accuracy ~97%
- 6. Prediction
 - Displayed predictions on test digits with matplotlib

Results:

Metric	Value
Test Accuracy	~97% (may vary)
Input Size	28 x 28
Output Labels	0 to 9

Conclusion:

The neural network model accurately classified handwritten digits using the MNIST dataset. With basic layers and ReLU activation, the model achieved strong performance. This demonstrates the power of deep learning in image classification and forms the foundation for advanced computer vision tasks.

In []: